


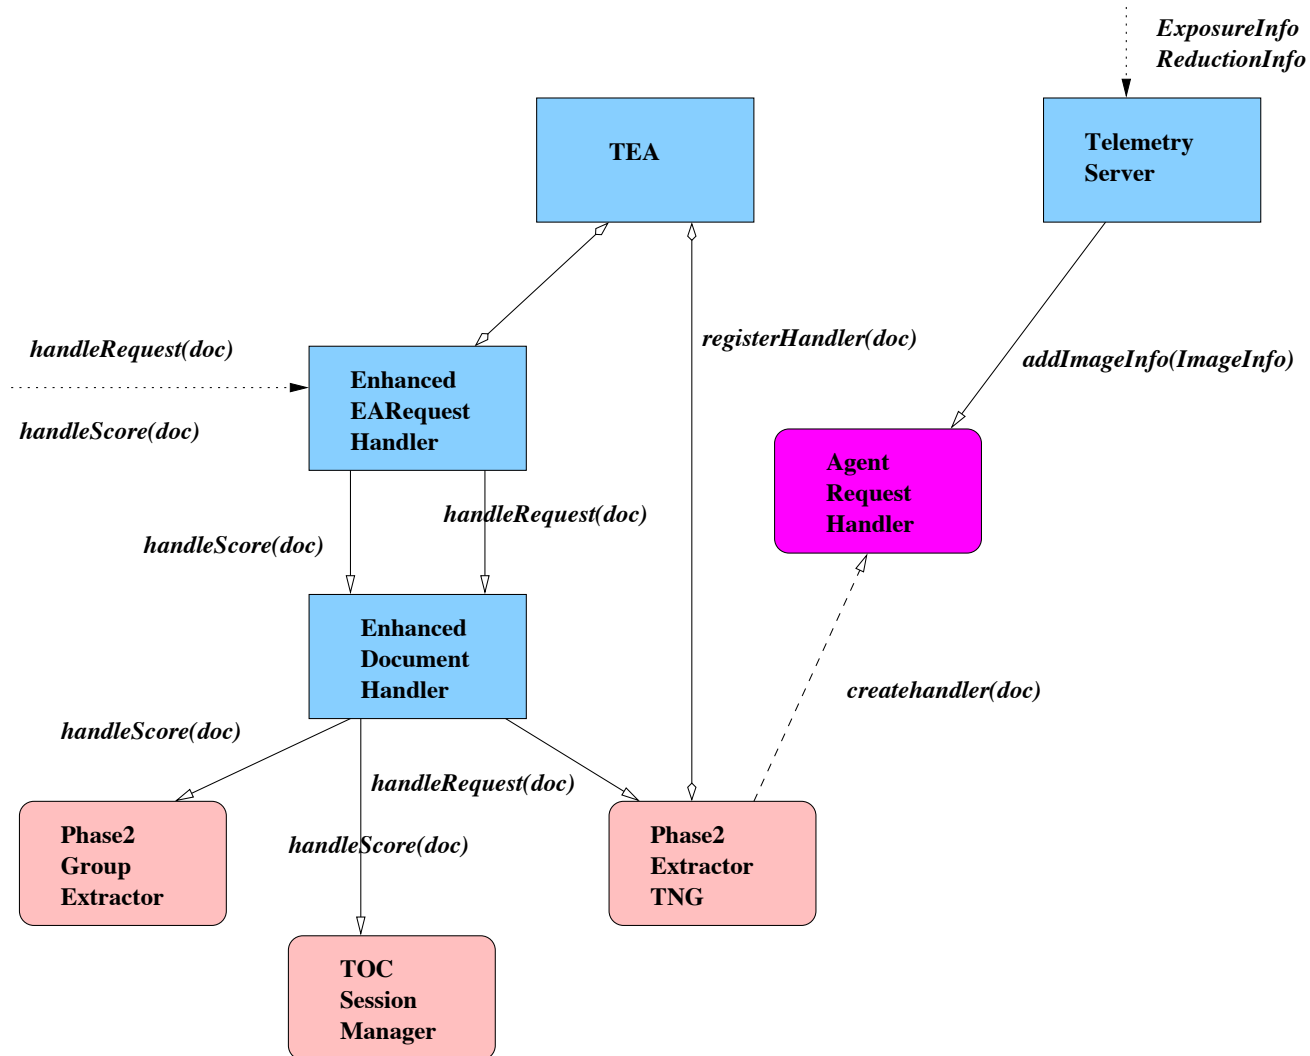
RTML AND NSO INTERFACES

TELESCOPE EMBEDDED AGENT

- The TEA was written several years ago and has gone through several iterations.
- Originally used estarIO as comms layer.
- Now uses RTML requests from the NodeAgent.
- Requires potentially rapid telemetry from the RCS to make asynch callbacks to remote agents.
- At the time of writing there was no suitable pub-sub telemetry system.
- Existing ports/holes were used to make a fudged solution – CAMP.
- Code is in *~dev/src/estar/tea/* and is under svn or maybe rcs ?




ARCHITECTURE




OPERATION

- Requests from the NA are passed to the *EmbeddedAgentRequestHandler* via calls to:-
 - `handleScore(RtmlDoc)` - scoring
 - `handleRequest(RtmlDoc)` – request to observe
- The observation, timing and constraint information is extracted and either a scoring decision or observation request may be made.
- TO requests are handled by
 - *TOCSessionManager*.
- Scheduled requests are handled by
 - `Phase2ExtractorTNG`.
- In the case of a scheduled request, an observation request is created and entered into the Phase2 ODB via the OSS. (rmi calls to `Phase2Model` etc).
- An *AgentRequestHandler* (ARQ) thread is started for each document currently in the Phase2 ODB to monitor for execution and make callbacks to the remote agent.
- Updates to the ARQ are received by the *TelemetryServer* in the form of:
 - `ObservationInfo` (a group has started)
 - `ObservationCompletionInfo` (a group has completed)
 - `ExposureInfo` (an image was taken)
 - `ReductionInfo` (a reduced image is ready)
- ARQ handlers are persisted over restarts by a simple serialization mechanism.




PROBLEMS

- The TEA is not very efficient – it was originally written to handle a few observation requests.
- The ARQs are separate threads – it would be better to have a single thread handle all the documents in round-robin fashion – much more scalable.
- The Telemetry mechanism is crude – there are modern feeds which could do this job anyway - OPS feed.
- The TEA keeps a local cache of proposals, groups and instrument-configs – this could be done nowadays via the Phase2 cache via rmi.
- RTML is structured like the old phase2 system.
- It is therefore necessary to translate from groups of observations to sequences and a certain amount of deduction is needed.



NSO INTERFACE

- This is an XML based interface.
- Phase 2 markup language (P2ML).
- Version (1) – based on the old phase2 with observations inside groups.
 - Most NSO observations still use this model.
 - They have to be translated into new sequences and constraint types.
- Version (2) – based on modern sequences.
 - Some NSO observations use this.
- All code is in the misc-systems repository and is in the default package.
 - V1 uses a SAX parser – it is evil. Implemented via *XMLParser* class.
 - V2 uses JDOM – it is great. Implemented via *P2mlParser* class



OPERATION

- Something on the NSO site generates XML documents. These are read by a client and sent via TCP to a server on *ltproxy*.
- Offline Control Relay (OCR) – cant recall where the name came from.
- Started using: */etc/init.d/ocr_init start*
- All the code is deployed to
 - */proxy/misc/misc_systems.jar*.
- Scripts are in
 - */proxy/bin/*
- Like the RTML interface – V1 requests need to be translated to sequences – a fair amount of deduction is needed.

